

Untitled

Ivo Pinheiro

2024-09-06

Practical Machine Learning Course Project

To predict the manner in which participants performed barbell lifts using sensor data, that was the task at hand. To accomplish this, I cleaned and analysed data to build a model that predicts the value of a target variable (\$classe) based on input variables (53 features).

```
# Libraries
set.seed(123)
library(ggplot2)
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##      margin
library(caret)

## Loading required package: lattice
```

Data preprocessing

I started by assigning NA to any missing value in the data, and then I removed any column that had over 90% NA. As it didn't make sense to include the first six columns, these were removed. Lastly, the target variable (\$classe) was converted into a factor variable.

These steps reduced the number of columns from 160 to 54, and ensured that our data was now ready for analysis and building our model.

```
# Importing
url_training <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
url_testing <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
# Reading
training_data <- read.csv(url_training)
testing_data <- read.csv(url_testing)
# Cleaning training set
dim(training_data)

## [1] 19622 160

training_data[training_data == ""] <- NA
proportion_NA_train <- colSums(is.na(training_data))/nrow(training_data)
```

```
training_data1 <- training_data[, proportion_NA_train < 0.9]
training_data2 <- training_data1[, -c(1:6)]
training_data2$classe <- as.factor(training_data2$classe)
```

Model building and evaluation

Since this is a classification problem and we already know the target variable (\$classe), the model that made the most sense to me was random forests. I trained the model on 70% of the training set and then I tested it on the other 30%.

```
# Model with training set
in_Train <- createDataPartition(y=training_data2$classe, p=0.7, list=FALSE)
trainingset <- training_data2[in_Train, ]
testingset <- training_data2[-in_Train, ]
model_1 <- randomForest(classe ~ ., data = trainingset)
print(model_1)
```

```
##
## Call:
## randomForest(formula = classe ~ ., data = trainingset)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 7
##
## OOB estimate of error rate: 0.31%
## Confusion matrix:
##      A    B    C    D    E class.error
## A 3905     0     0     0     1 0.0002560164
## B   7 2648     3     0     0 0.0037622272
## C   0   10 2386     0     0 0.0041736227
## D   0    0  14 2237     1 0.0066607460
## E   0    0    0   6 2519 0.0023762376
```

Happy with the results, I applied the model to the testing set.

```
# Cleaning test set
dim(testing_data)
```

```
## [1] 20 160
```

```
testing_data[testing_data == ""] <- NA
proportion_NA_test <- colSums(is.na(testing_data))/nrow(testing_data)
testing_data1 <- testing_data[, proportion_NA_test < 0.9]
testing_data2 <- testing_data1[, -c(1:6)]
# Applying model to test set
testset_predictions <- predict(model_1, newdata = testing_data2)
print(testset_predictions)
```

```
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```

So what is the percentage of the target variable in the test set that are correctly classified by the model?

```
predictions <- predict(model_1, testingset)
confusionMatrix(predictions, testingset$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1674    0    0    0    0
##           B    0 1137    1    0    0
##           C    0    2 1025    2    0
##           D    0    0    0  962    4
##           E    0    0    0    0 1078
##
## Overall Statistics
##
##           Accuracy : 0.9985
##           95% CI : (0.9971, 0.9993)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9981
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9982  0.9990  0.9979  0.9963
## Specificity      1.0000  0.9998  0.9992  0.9992  1.0000
## Pos Pred Value   1.0000  0.9991  0.9961  0.9959  1.0000
## Neg Pred Value    1.0000  0.9996  0.9998  0.9996  0.9992
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2845  0.1932  0.1742  0.1635  0.1832
## Detection Prevalence 0.2845  0.1934  0.1749  0.1641  0.1832
## Balanced Accuracy 1.0000  0.9990  0.9991  0.9986  0.9982
```

According to the confusion matrix, the model correctly classified 99.85% of the instances overall.


Given the high accuracy, the expected out-of-sample error is likely to be very low ($1 - 0.9985 = 0.0015$).

Thus, the expected out-of-sample error is approximately 0.15%.

Plots

```
varImpPlot(model_1)
```

model_1



```
ggplot(training_data2, aes(x = classe)) +
  geom_bar(fill = "white", color = "grey") +
  labs(title = "How many of each $classe in the training set", x = NULL, y = NULL)
```

How many of each \$classe in the training set

